

# Apache Kafka

## Scalable Message Processing and more!

Guido Schmutz - 24.4.2017



@gschmutz



guidoschmutz.wordpress.com

BASEL ▪ BERN ▪ BRUGG ▪ DÜSSELDORF ▪ FRANKFURT A.M. ▪ FREIBURG I.BR. ▪ GENF  
HAMBURG ▪ KOPENHAGEN ▪ LAUSANNE ▪ MÜNCHEN ▪ STUTTGART ▪ WIEN ▪ ZÜRICH

**trivadis**  
makes **IT** easier. ■ ■ ■

# ■ Guido Schmutz

Working at Trivadis for more than 20 years

Oracle ACE Director for Fusion Middleware and SOA



Consultant, Trainer Software Architect for Java, Oracle, SOA and Big Data / Fast Data

Member of Trivadis Architecture Board

Technology Manager @ Trivadis



More than 30 years of software development experience

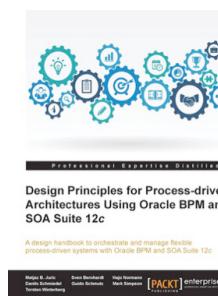
Contact: [guido.schmutz@trivadis.com](mailto:guido.schmutz@trivadis.com)

Blog: <http://guidoschmutz.wordpress.com>

Slideshare: <http://www.slideshare.net/gschmutz>

Twitter: [@gschmutz](https://twitter.com/gschmutz)

Apache Kafka - Scalable Message Processing and more!



**trivadis**  
makes IT easier.

# ■ Agenda

1. Introduction & Motivation
2. Kafka Core
3. Kafka Connect
4. Kafka Streams
5. Kafka and "Big Data" / "Fast Data" Ecosystem
6. Kafka in Enterprise Architecture
7. Confluent Data Platform
8. Summary

Apache Kafka - Scalable Message Processing and more!

**trivadis**  
makes **IT** easier. 

# Introduction & Motivation

Apache Kafka - Scalable Message Processing and more!

**trivadis**  
makes **IT** easier. 

## ■ Apache Kafka - Overview



Distributed publish-subscribe messaging system

Designed for processing of real time activity stream data (logs, metrics collections, social media streams, ...)

Initially developed at LinkedIn, now part of Apache

Does not use JMS API and standards

Kafka maintains feeds of messages in topics

## ■ Apache Kafka - Motivation



LinkedIn's motivation for Kafka was:

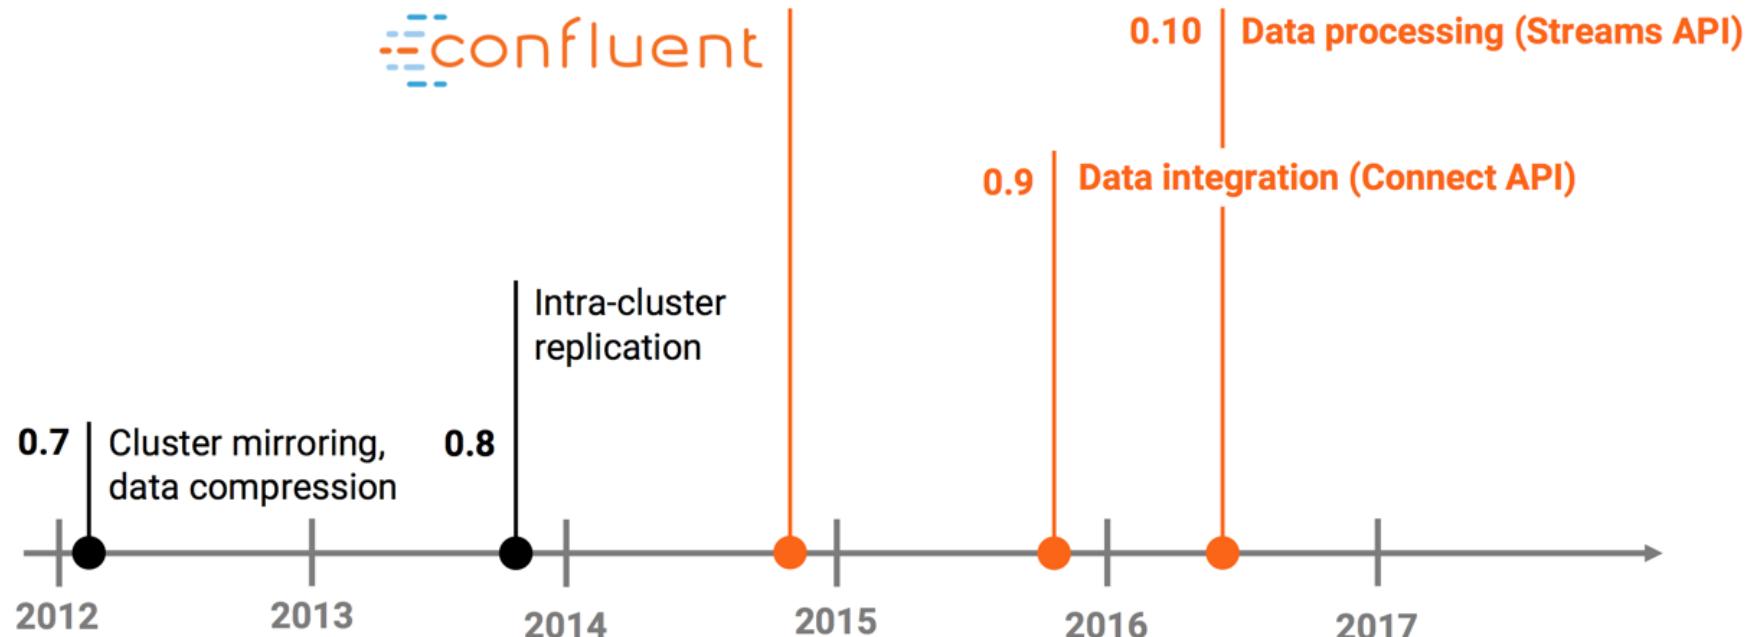
- "A unified platform for handling all the real-time data feeds a large company might have."

Must haves

- High throughput to support **high volume event feeds**
- Support real-time processing of these feeds to create **new, derived feeds**.
- Support large data backlogs to handle periodic ingestion from **offline systems**
- Support low-latency delivery to handle more traditional **messaging use cases**
- Guarantee **fault-tolerance** in the presence of machine failures

# ■ Apache Kafka History

Apache Kafka: birthed as a messaging system, now a **streaming platform**

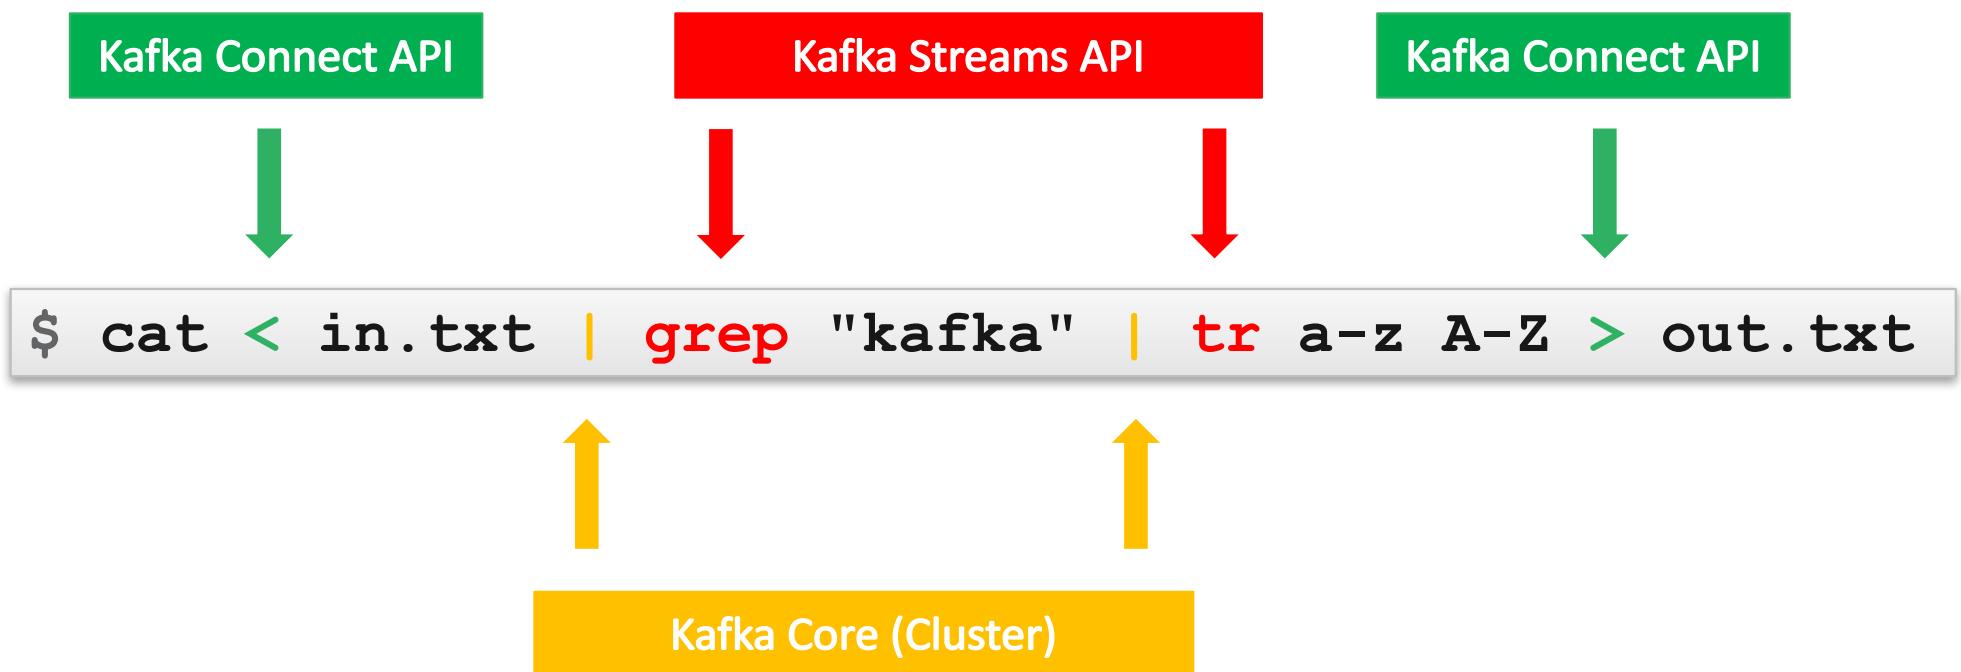


Source: Confluent

Apache Kafka - Scalable Message Processing and more!

**trivadis**  
makes IT easier. ■ ■ ■

## ■ Apache Kafka - Unix Analogy



Source: Confluent

Apache Kafka - Scalable Message Processing and more!

**trivadis**  
makes **IT** easier. . .

# Kafka Core

Apache Kafka - Scalable Message Processing and more!

**trivadis**  
makes **IT** easier. 

# ■ Kafka High Level Architecture

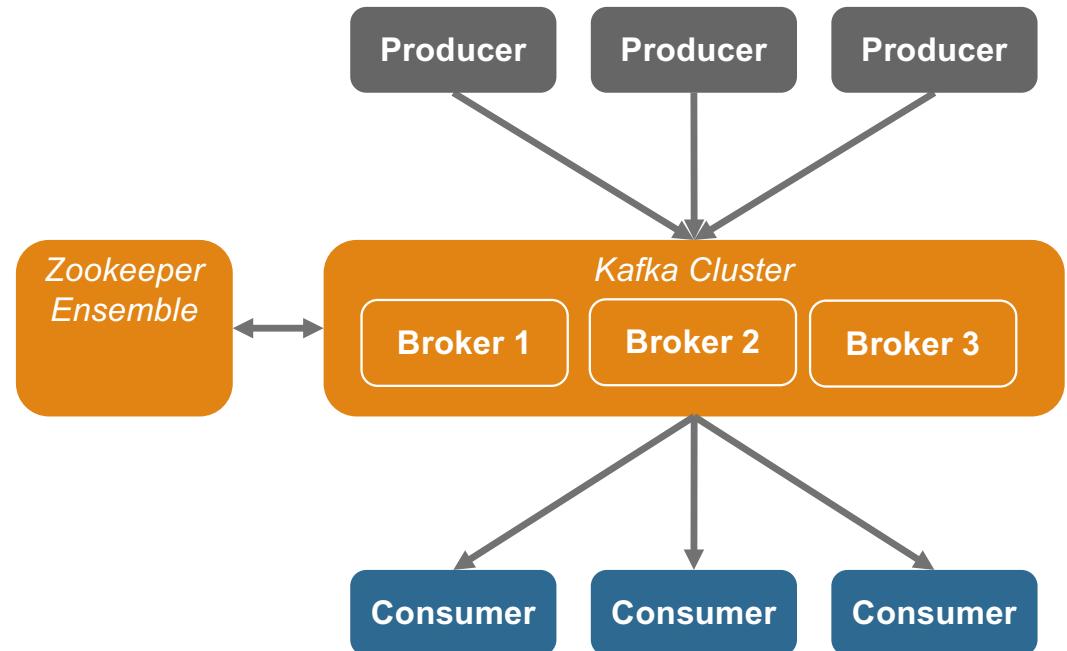


The who is who

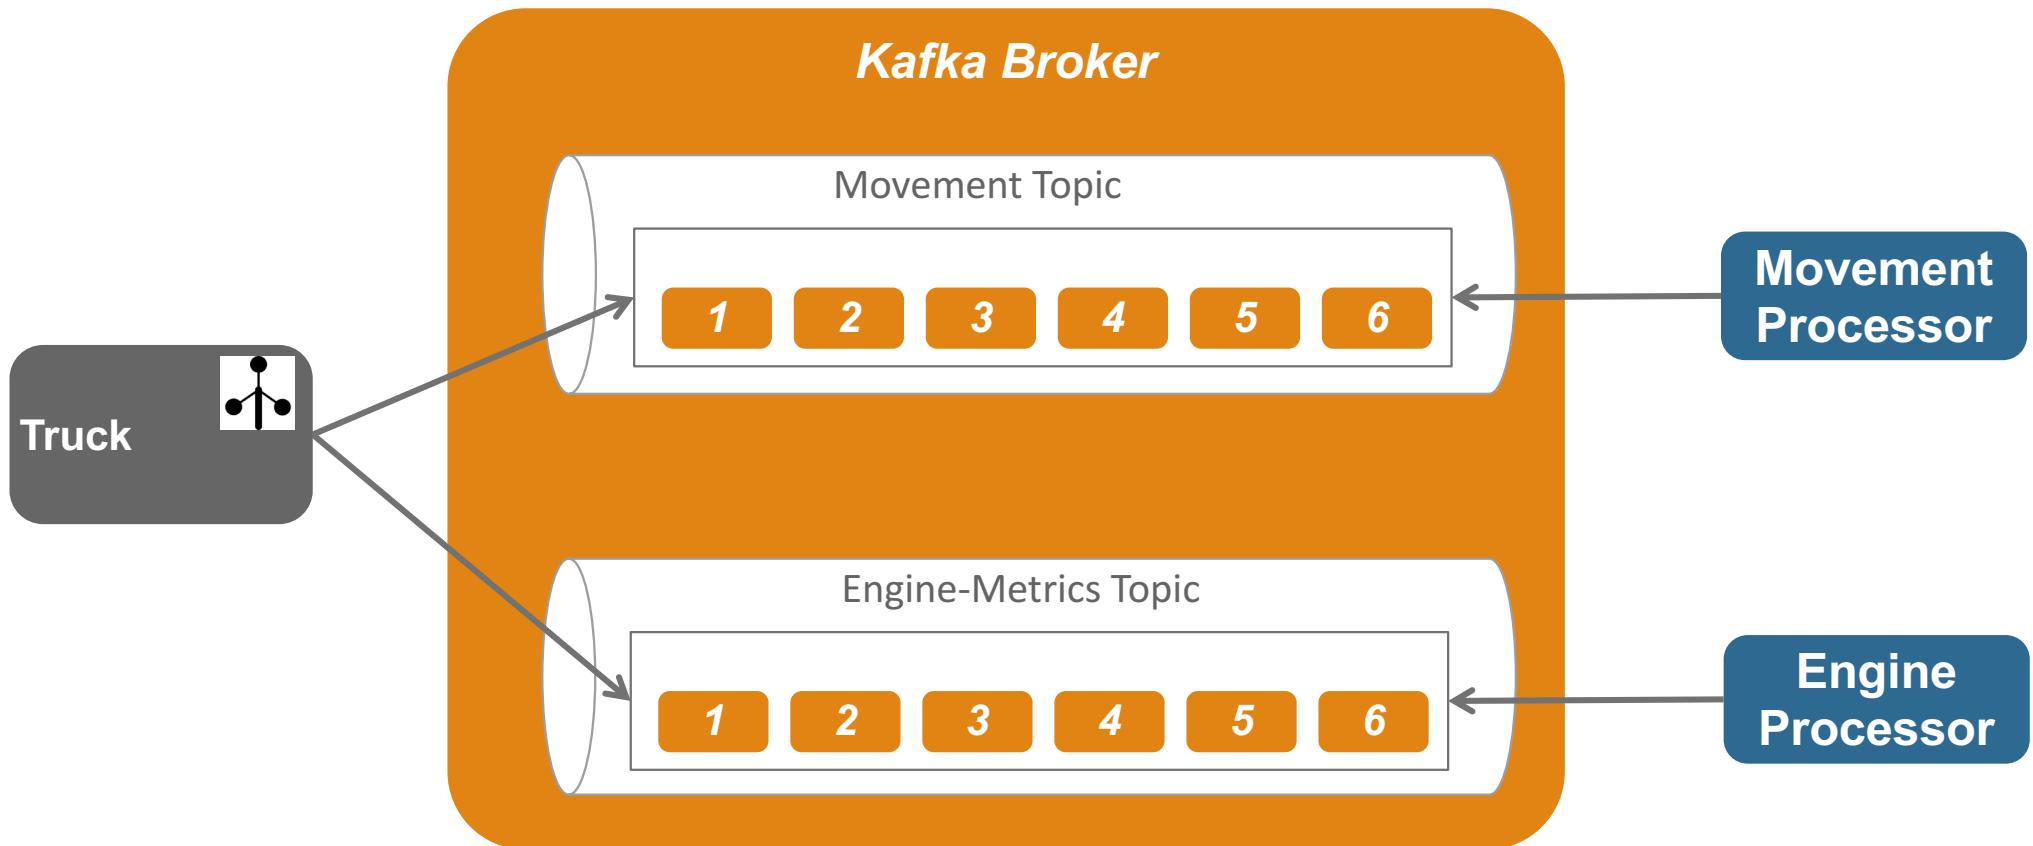
- **Producers** write data to **brokers**.
- **Consumers** read data from **brokers**.
- All this is distributed.

The data

- Data is stored in **topics**.
- **Topics** are split into **partitions**, which are **replicated**.



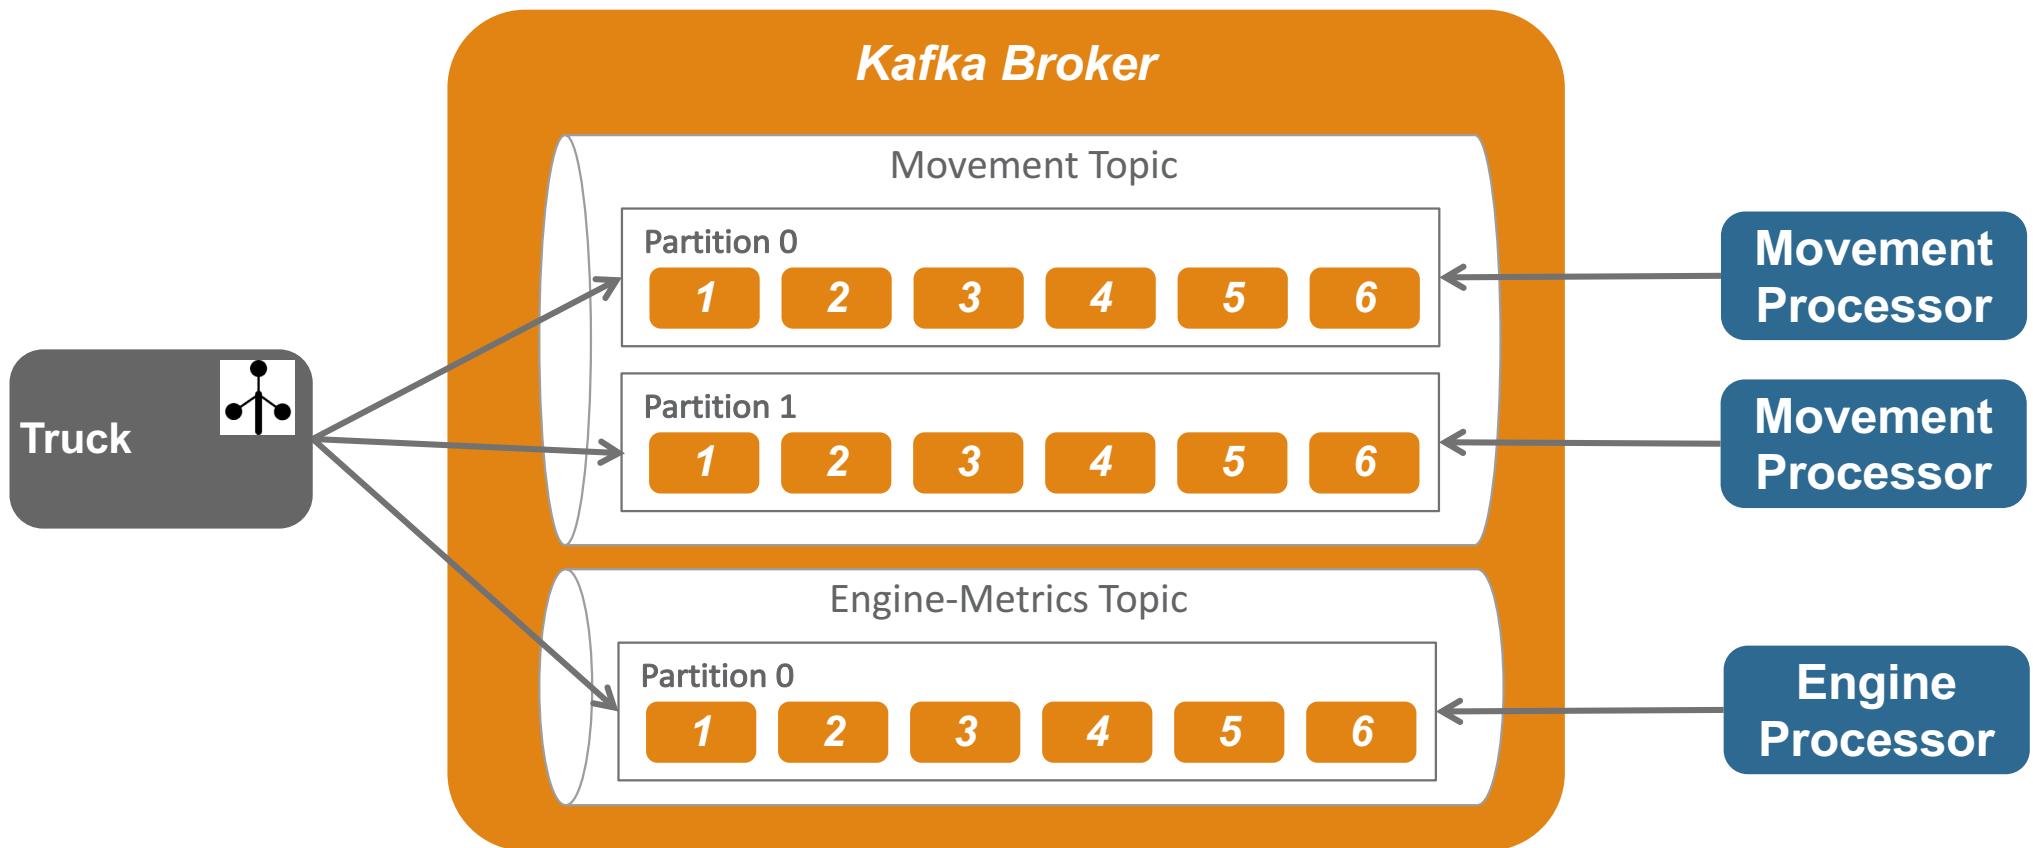
# ■ Apache Kafka - Architecture



Apache Kafka - Scalable Message Processing and more!

**trivadis**  
makes IT easier. ■ ■ ■

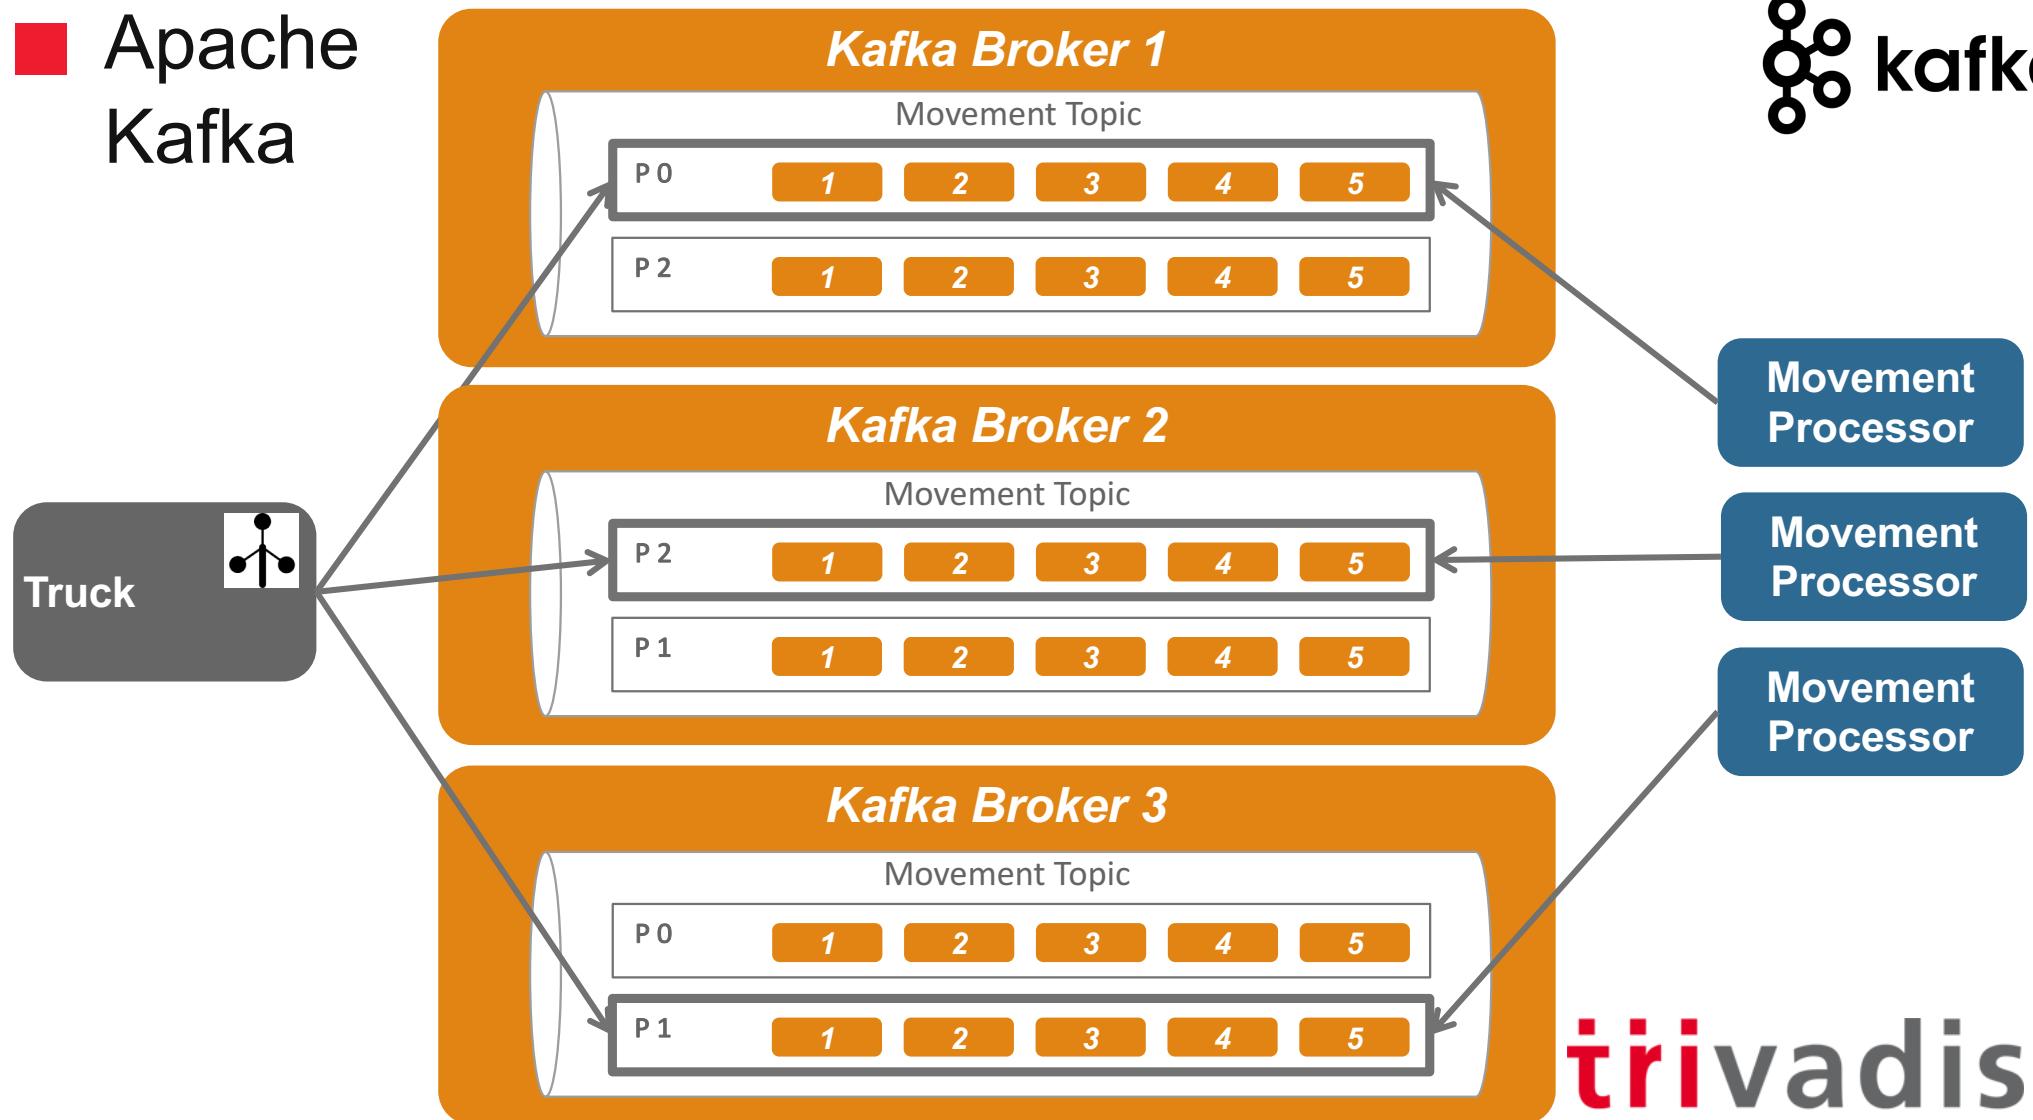
# ■ Apache Kafka - Architecture



Apache Kafka - Scalable Message Processing and more!

**trivadis**  
makes IT easier. ■ ■ ■

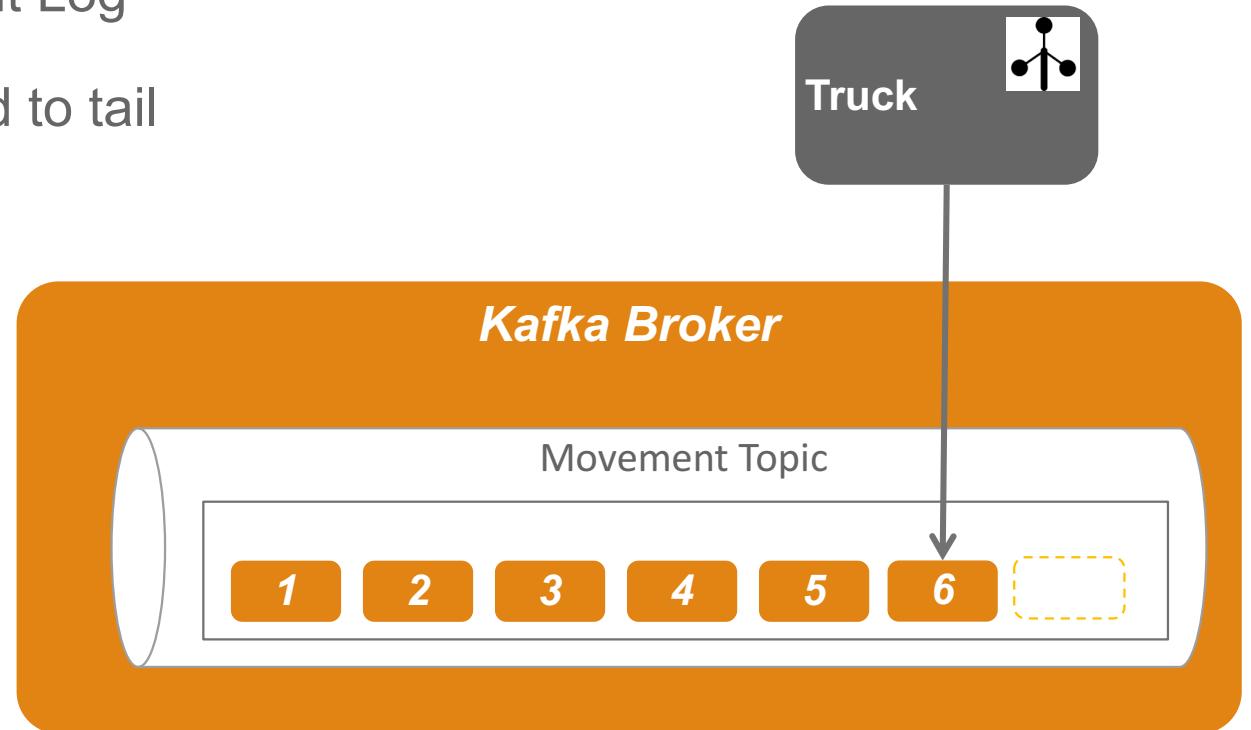
# ■ Apache Kafka



## ■ Apache Kafka - Architecture



- Write Ahead Log / Commit Log
- Producers always append to tail
- think append to file



# ■ Kafka Topics

## Creating a topic

- Command line interface

```
$ kafka-topics.sh --zookeeper zk1:2181 --create \
    --topic my.topic --partitions 3 \
    --replication-factor 2 --config x=y
```

- Using AdminUtils.createTopic method
- Auto-create via auto.create.topics.enable = true

## Modifying a topic

[https://kafka.apache.org/documentation.html#basic\\_ops\\_modify\\_topic](https://kafka.apache.org/documentation.html#basic_ops_modify_topic)

## Deleting a topic

- Command Line interface

## ■ Kafka Producer

```
private Properties kafkaProps = new Properties();
kafkaProps.put("bootstrap.servers", "broker1:9092,broker2:9092");
kafkaProps.put("key.serializer", "...StringSerializer");
kafkaProps.put("value.serializer", "...StringSerializer");

producer = new KafkaProducer<String, String>(kafkaProps);
```

```
ProducerRecord<String, String> record =
    new ProducerRecord<>("topicName", "Key", "Value");
try {
    producer.send(record);
} catch (Exception e) {}
```

# Durability Guarantees

Producer can configure acknowledgements

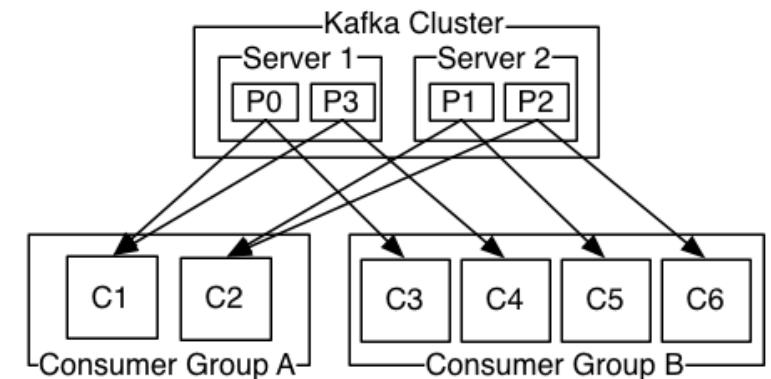
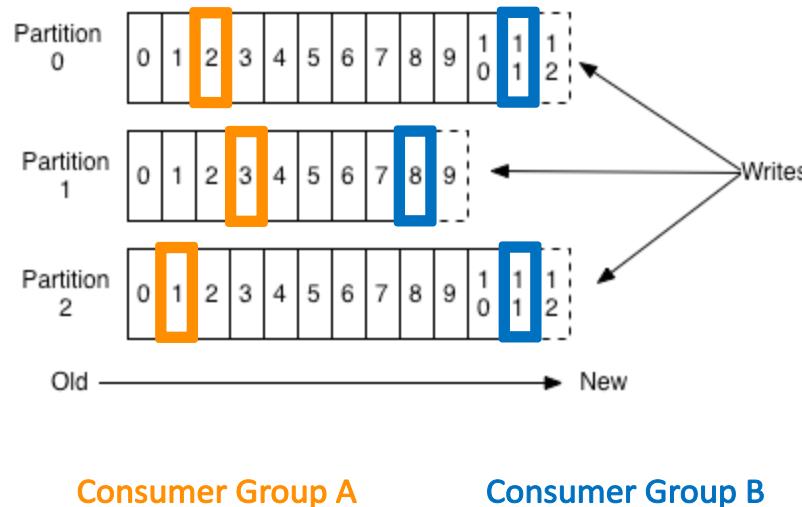
Value	Description	Throughput	Latency	Durability
0	<ul style="list-style-type: none"><li>Producer doesn't wait for leader</li></ul>	high	low	low (no guarantee)
1 (default)	<ul style="list-style-type: none"><li>Producer waits for leader</li><li>Leader sends ack when message written to log</li><li>No wait for followers</li></ul>	medium	medium	medium (leader)
all (-1)	<ul style="list-style-type: none"><li>Producer waits for leader</li><li>Leader sends ack when all In-Sync Replica have acknowledged</li></ul>	low	high	high (ISR)

# ■ Apache Kafka - Partition offsets



Offset: messages in the partitions are each assigned a unique (per partition) and sequential id called the offset

- Consumers track their pointers via *(offset, partition, topic)* tuples



Source: Apache Kafka

## ■ Data Retention – 3 options

1. Never
2. Time based (TTL)

```
log.retention.{ms | minutes | hours}
```

3. Size based

```
log.retention.bytes
```

4. Log compaction based (entries with same key are removed)

```
kafka-topics.sh --zookeeper localhost:2181 \
    --create --topic customers \
    --replication-factor 1 --partitions 1 \
    --config cleanup.policy=compact
```

## ■ Apache Kafka – Some numbers



Kafka at LinkedIn => over 1800+ broker machines / 79K+ Topics

<https://engineering.linkedin.com/kafka/running-kafka-scale>

1.3 Trillion messages per day

330 Terabytes in/day  
1.2 Petabytes out/day

Peak load for a single cluster  
2 million messages/sec  
4.7 Gigabits/sec inbound  
15 Gigabits/sec outbound

Kafka Performance at our own infrastructure => 6 brokers (VM) / 1 cluster

<http://engineering.linkedin.com/kafka/benchmarking-apache-kafka-2-million-writes-second-three-cheap-machines>

- 445'622 messages/second
- 31 MB / second
- 3.0405 ms average latency between producer / consumer

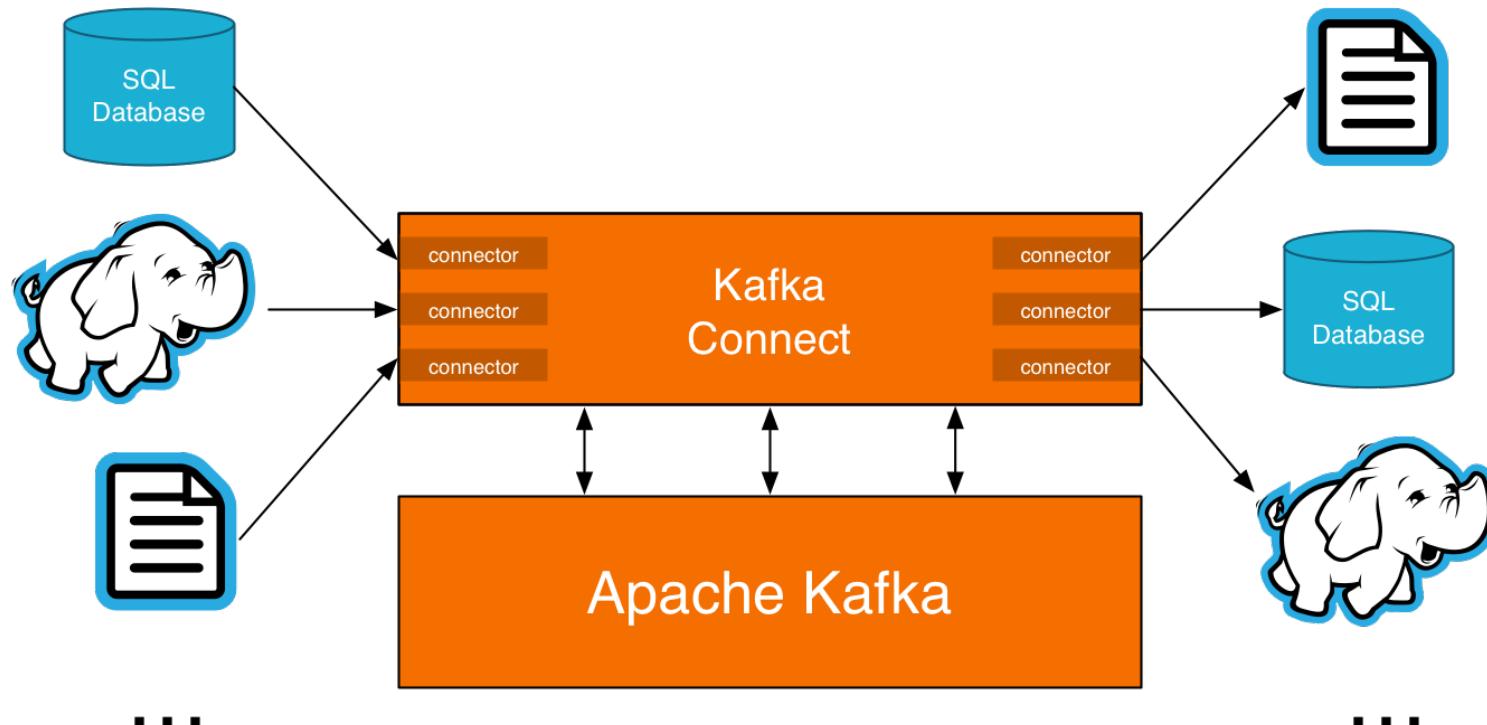
# Kafka Connect

Apache Kafka - Scalable Message Processing and more!

**trivadis**  
makes **IT** easier. 

# Kafka Connect Architecture

Data Sources    Data Sinks



Source: Confluent

Apache Kafka - Scalable Message Processing and more!

**trivadis**  
makes IT easier. ■ ■ ■

# Kafka Connector Hub – Certified Connectors

CONNECTOR	TAGS	DEVELOPER	SUPPORT
HDFS (Sink)	HDFS, Hadoop, Hive	Confluent	Confluent
JDBC (Source)	JDBC, MySQL	Confluent	Confluent
Attunity (Sink)	CDC, Oracle	Attunity	Attunity
Couchbase (Source)	Couchbase, NoSQL	Couchbase	Couchbase
JustOne (Sink)	Postgress	JustOne	JustOne
Striim (Source)	CDC, Oracle, MS SQLServer	Striim	Striim
Syncsort DMX (Source)	DB2, IMS, VSAM, CICS	Syncsort	Syncsort
Syncsort DMX (Sink)	DB2, IMS, VSAM, CICS	Syncsort	Syncsort
Vertica (Source)	Vertica	HP Enterprise	HP Enterprise
Vertica (Sink)	Vertica	HP Enterprise	HP Enterprise

Source: <http://www.confluent.io/product/connectors>

Apache Kafka - Scalable Message Processing and more!

**trivadis**  
makes **IT** easier. ■ ■ ■

# ■ Kafka Connector Hub – Additional Connectors

CONNECTOR	TAGS	DEVELOPER	SUPPORT
Apache Ignite (Source)	File System	Community	Community
Apache Ignite (Sink)	File System	Community	Community
Bloomberg Ticker (Source)		Community	Community
Cassandra (Source)	Cassandra, Datastax	Community	Community 1 Community 2
Cassandra (Sink)	Cassandra, Datastax	Community	Community 1 Community 2
Elastic Search (Sink)	search, elastic, log, analytics	Community	Community 1 Community 2 Community 3
HBase Sink	HBase, NoSQL	Community	Community
Kudu (Sink)	Kudu	Community	Community
Mixpanel (Source)	analytics	Community	Community
MongoDB (Source)	Mongo, MongoDB, NoSQL	Community	Community
MQTT (Source)	MQTT, messaging	Community	Community
MySQL CDC - Debezium (Source)	MySQL, CDC, Oracle	Community	Community

Source: <http://www.confluent.io/product/connectors>

Apache Kafka - Scalable Message Processing and more!

**trivadis**  
makes **IT** easier. ■ ■ ■

## ■ Kafka Connect – Twitter example

```
./connect-standalone.sh ../demo-config/connect-simple-source-standalone.properties  
..../demo-config/twitter-source.properties
```

```
bootstrap.servers=localhost:9095,localhost:9096,localhost:9097  
  
key.converter=org.apache.kafka.connect.storage.StringConverter  
value.converter=org.apache.kafka.connect.storage.StringConverter  
...
```

```
name=twitter-source  
connector.class=com.eneco.trading.kafka.connect.twitter.TwitterSourceConnector  
tasks.max=1  
topic=tweets  
twitter.consumerkey=<consumer-key>  
twitter.consumersecret=<consumer-secret>  
twitter.token=<token>  
twitter.secret=<token-secret>  
track.terms=bigdata
```

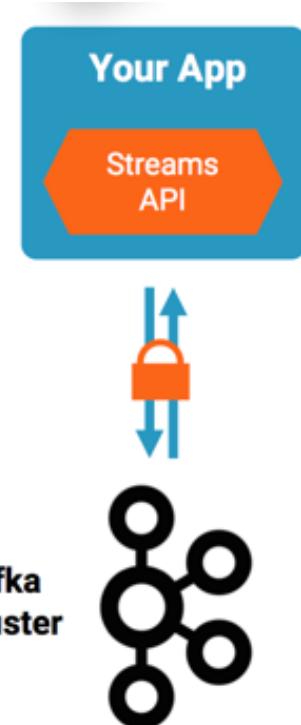
# Kafka Streams

Apache Kafka - Scalable Message Processing and more!

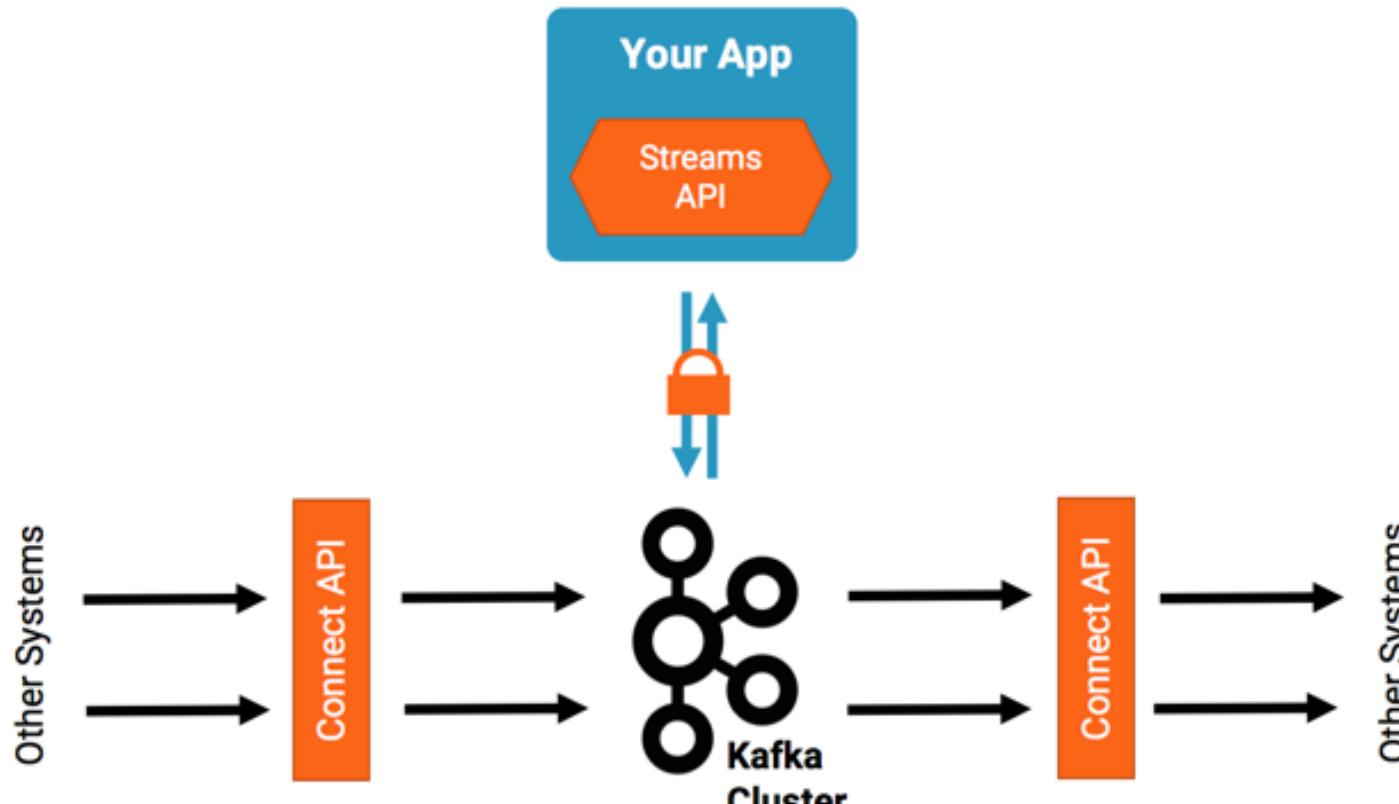
**trivadis**  
makes **IT** easier. 

# Kafka Streams

- Designed as a **simple and lightweight library** in Apache Kafka
- no external dependencies on systems other than Apache Kafka
- Part of open source Apache Kafka, introduced in 0.10+
- Leverages **Kafka as its internal messaging layer**
- agnostic to resource management and configuration tools
- Supports **fault-tolerant local state**
- Event-at-a-time processing (not microbatch) with millisecond latency
- Windowing with out-of-order data using a Google DataFlow-like model



## ■ Streams API in the context of Kafka



Source: Confluent

Apache Kafka - Scalable Message Processing and more!

**trivadis**  
makes **IT** easier. ■ ■ ■

# Kafka and "Big Data" / "Fast Data" Ecosystem

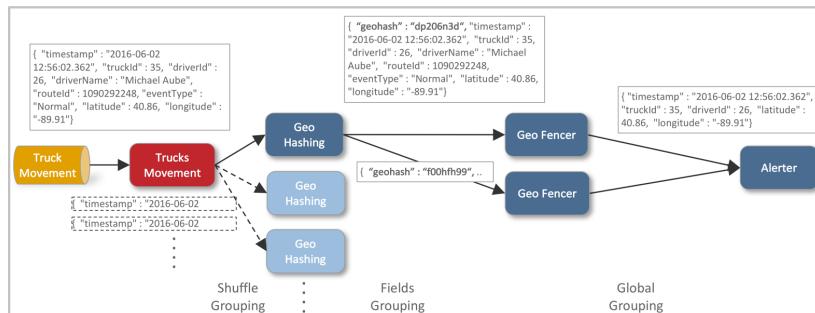
Apache Kafka - Scalable Message Processing and more!

**trivadis**  
makes **IT** easier. 

# ■ Kafka and the Big Data / Fast Data ecosystem

Kafka integrates with many popular products / frameworks

- Apache Spark Streaming
- Apache Flink
- Apache Storm



Storm built-in Kafka Spout to consume events from Kafka

- Apache NiFi
- Streamsets
- Apache Flume
- Oracle Stream Analytics
- Oracle Service Bus
- Oracle GoldenGate
- Spring Integration Kafka Support
- ...

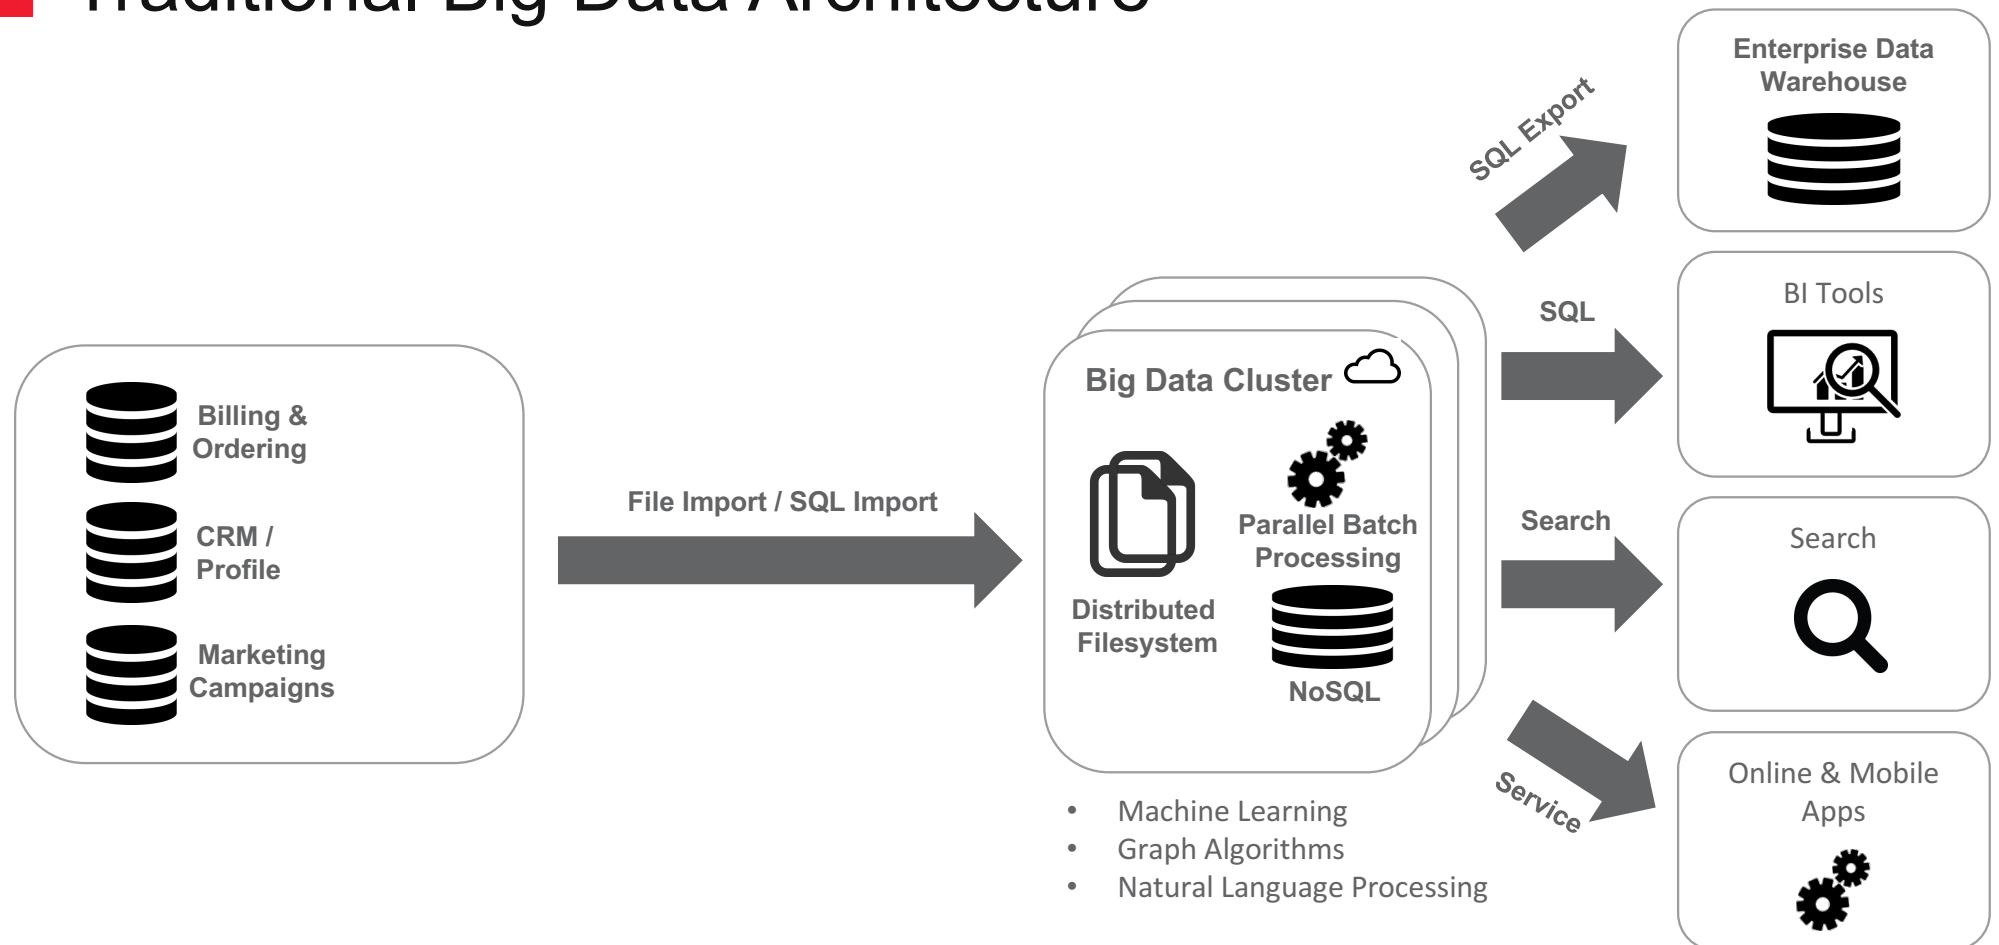
Apache Kafka - Scalable Message Processing and more!

# Kafka in “Enterprise Architecture”

Apache Kafka - Scalable Message Processing and more!

**trivadis**  
makes **IT** easier. 

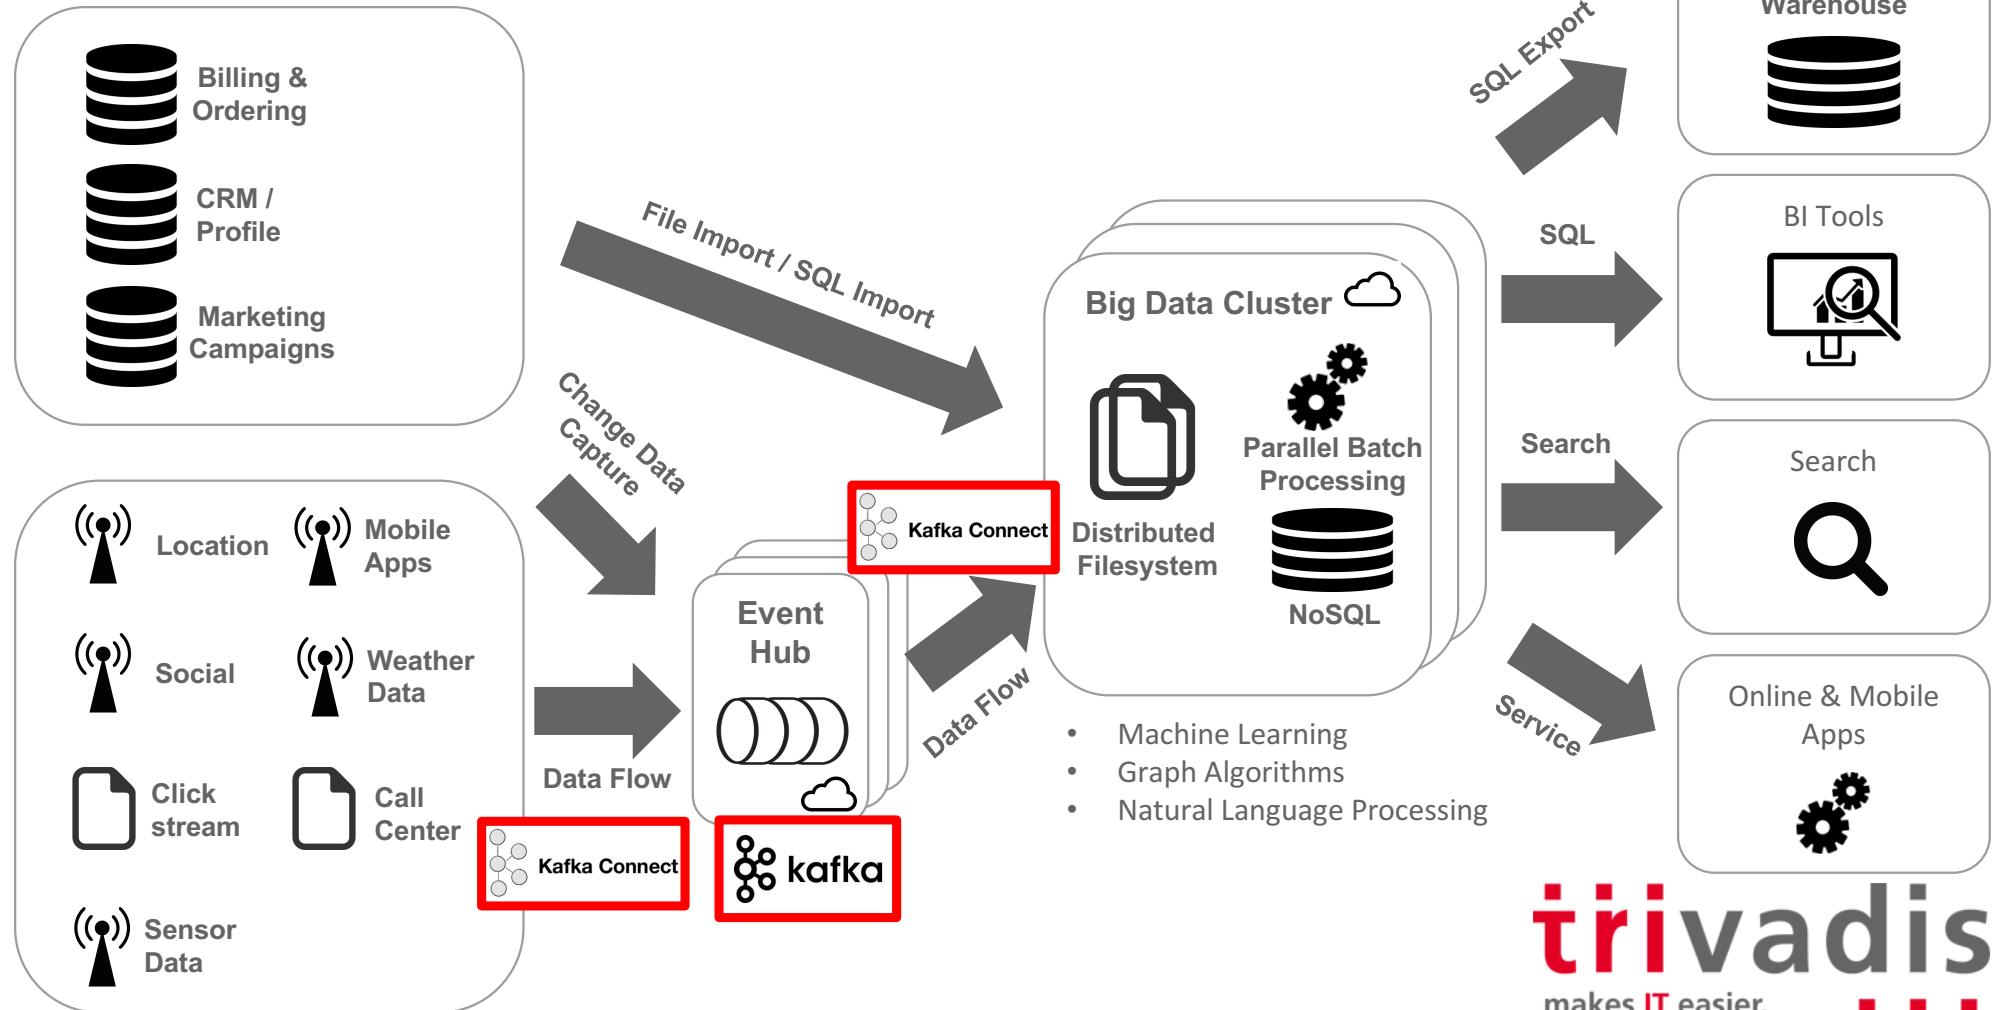
# ■ Traditional Big Data Architecture



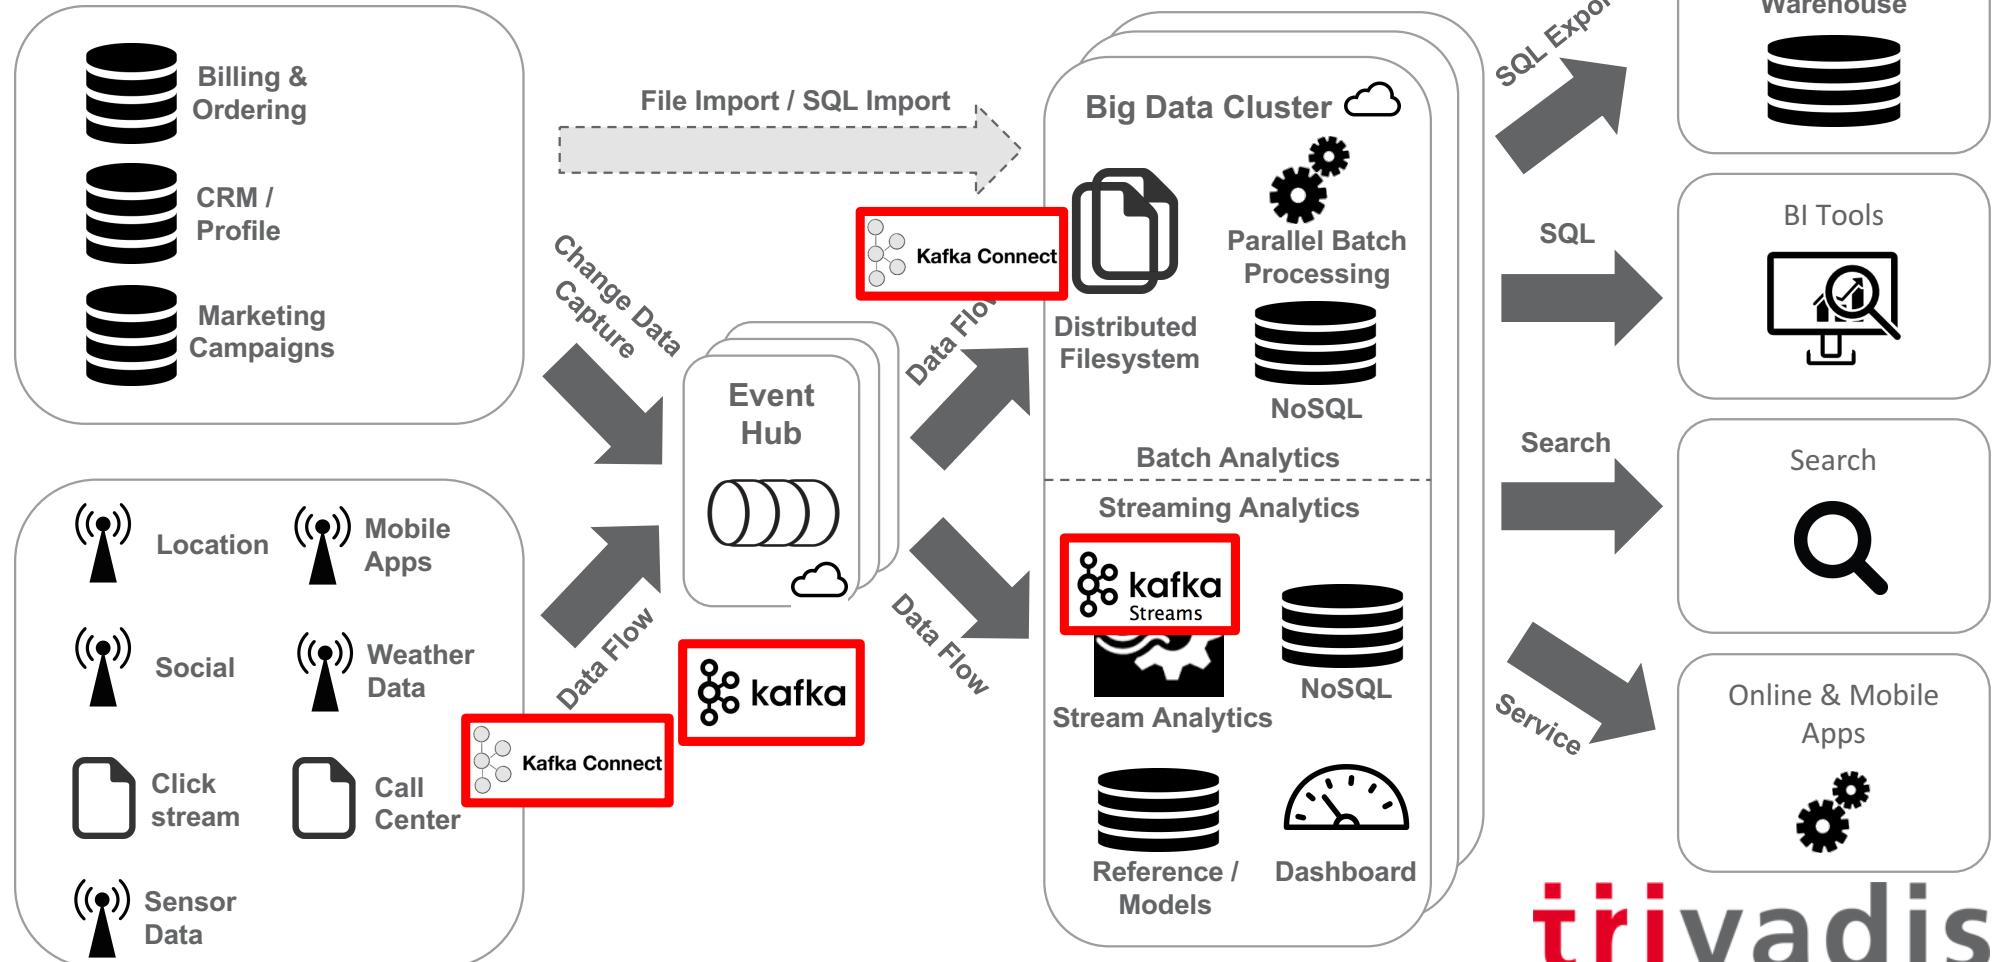
Apache Kafka - Scalable Message Processing and more!

**trivadis**  
makes **IT** easier. ■ ■ ■

## ■ Event Hub – handle event stream data



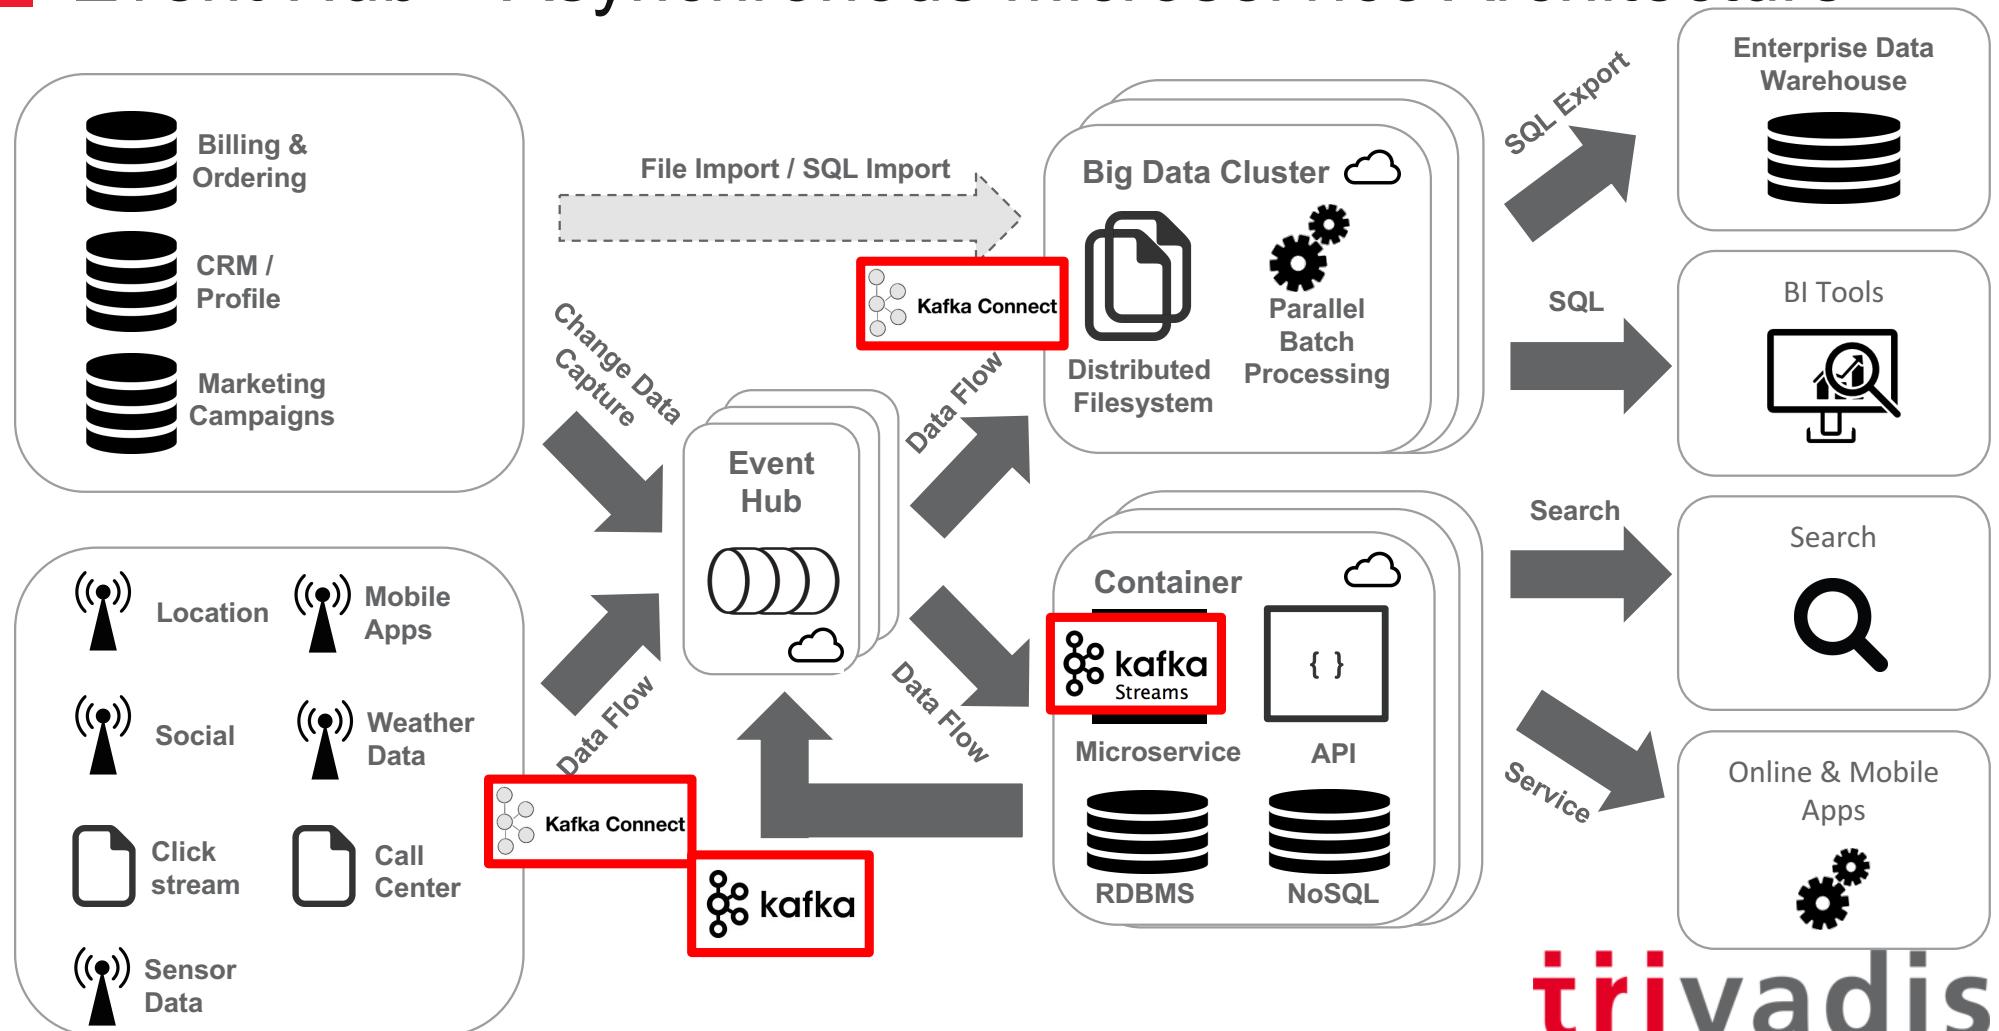
# ■ Event Hub – taking Velocity into account



Apache Kafka - Scalable Message Processing and more!

**trivadis**  
makes IT easier. ■ ■ ■

# Event Hub – Asynchronous Microservice Architecture



Apache Kafka - Scalable Message Processing and more!

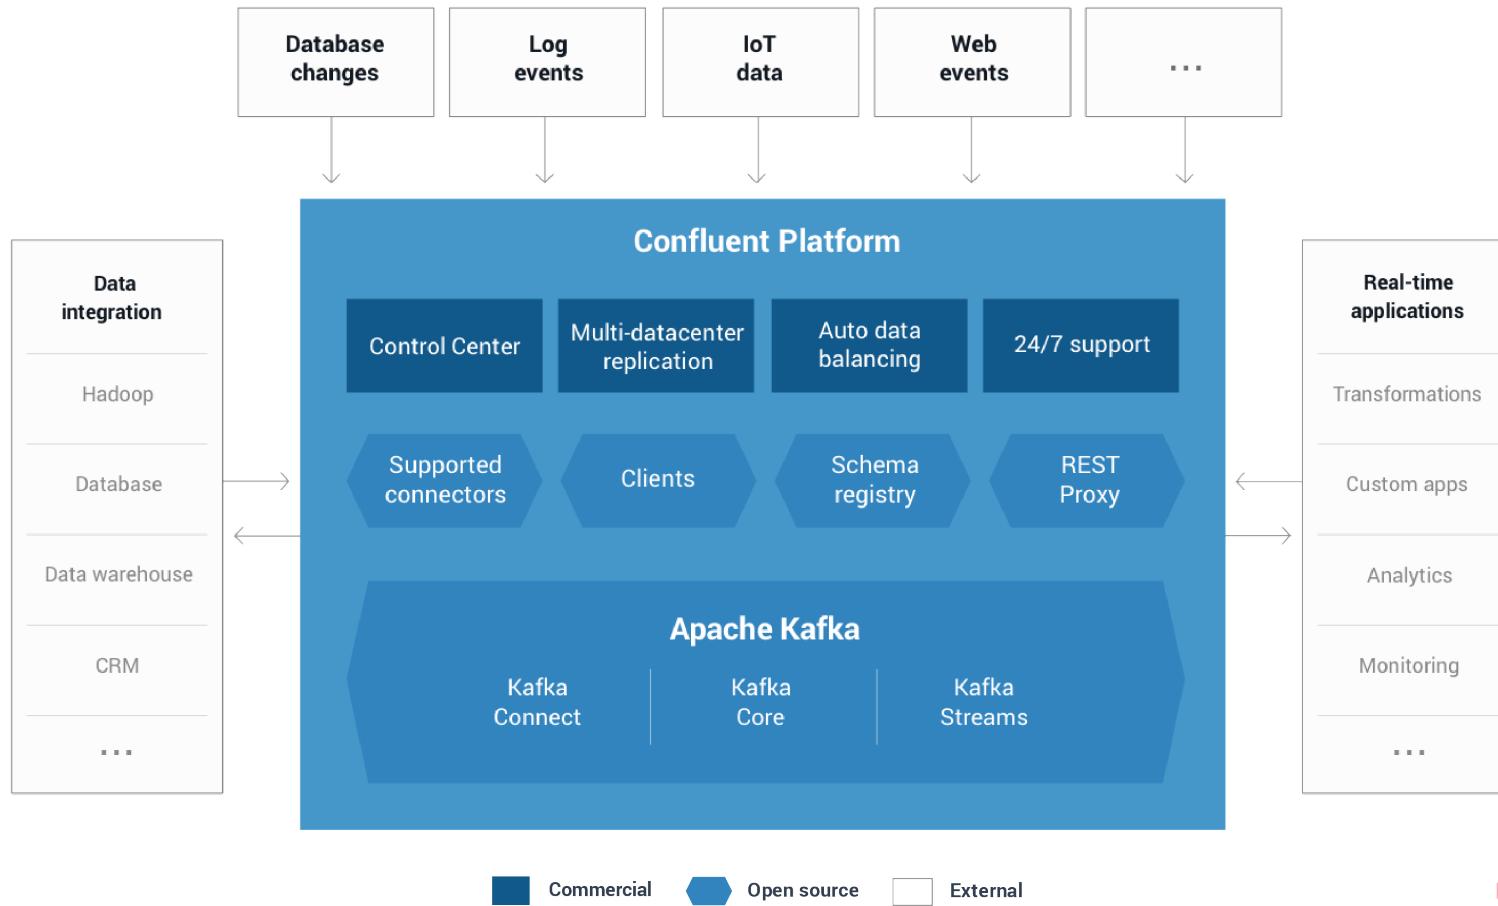
**trivadis**  
makes IT easier. ■ ■ ■

# Confluent Platform

Apache Kafka - Scalable Message Processing and more!

**trivadis**  
makes **IT** easier. 

# Confluent Data Platform 3.2



Source: Confluent

Apache Kafka - Scalable Message Processing and more!

**trivadis**  
makes IT easier. ■ ■ ■

# ■ Confluent Data Platform 3.2



## Confluent Open Source

Open source package

Connectors

Schema Registry

Clients

REST Proxy



## Confluent Enterprise

Enterprise package

Connectors

Schema Registry

Clients

REST Proxy

Control Center

Auto Data Balancing

Multi DC Replication

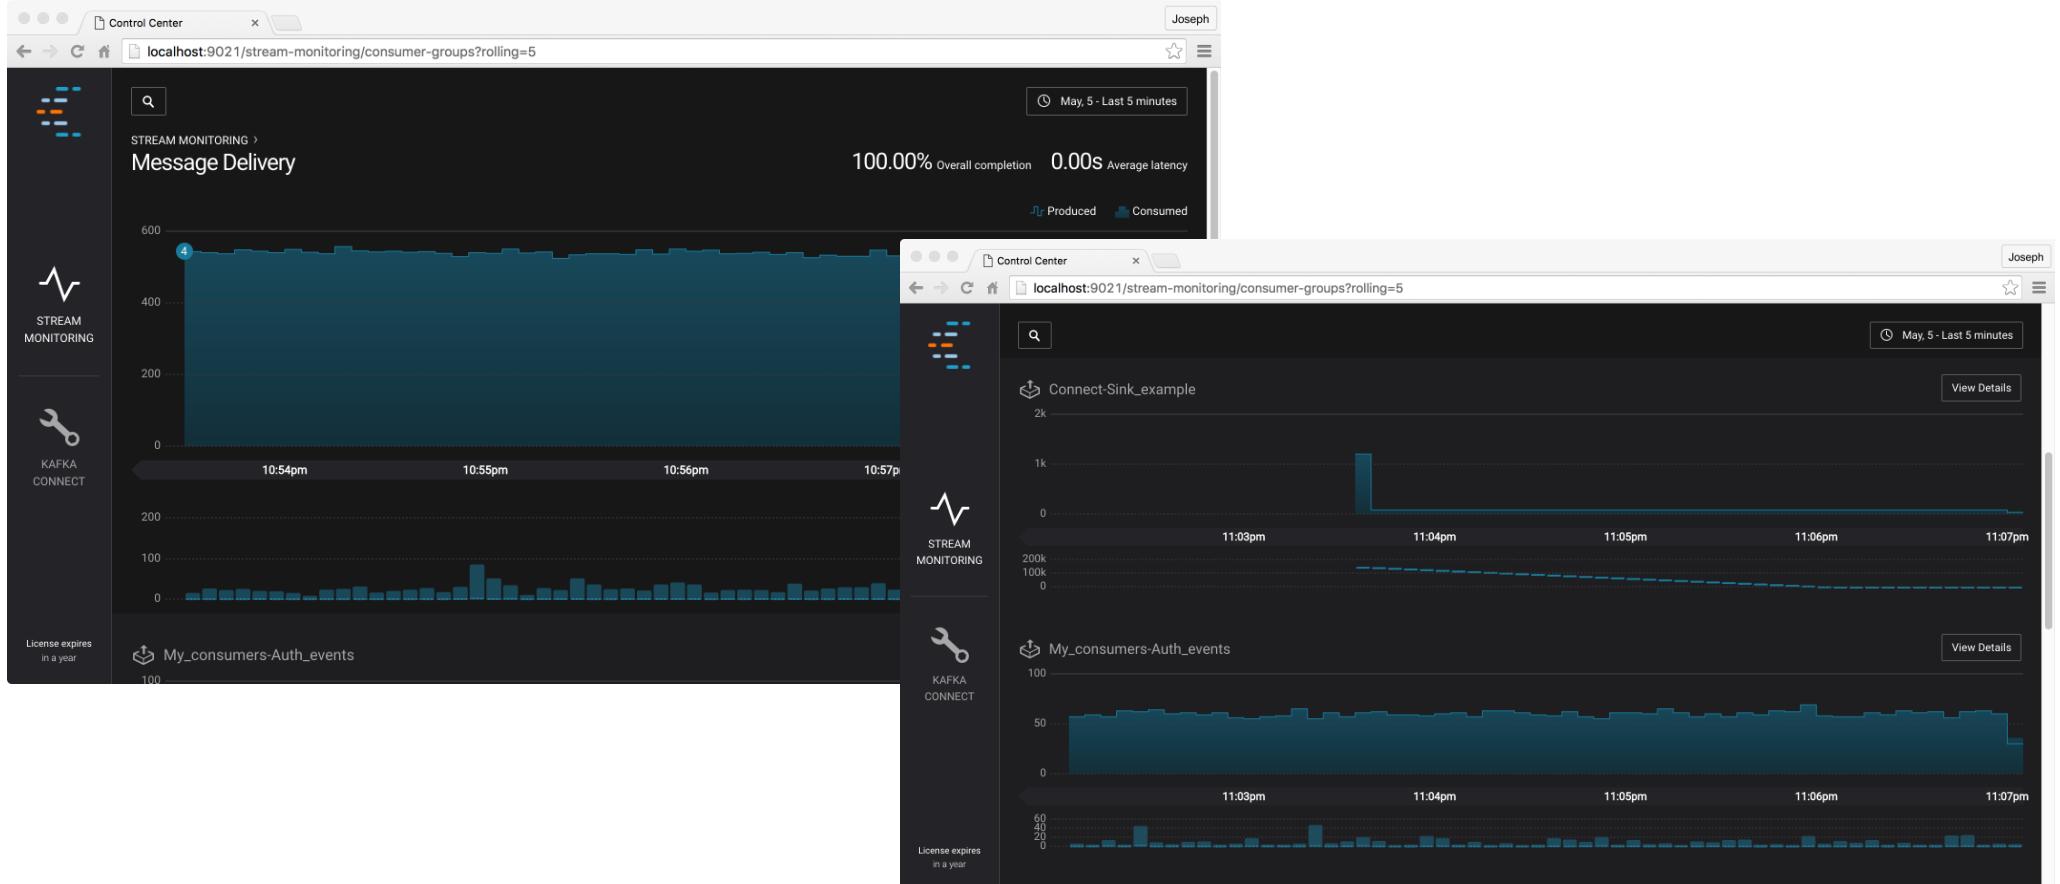
24/7 Support

Source: Confluent

Apache Kafka - Scalable Message Processing and more!

**trivadis**  
makes **IT** easier. ■ ■ ■

# Confluent Enterprise – Control Center



Source: Confluent

Apache Kafka - Scalable Message Processing and more!

**trivadis**  
makes IT easier. ■ ■ ■

# Summary

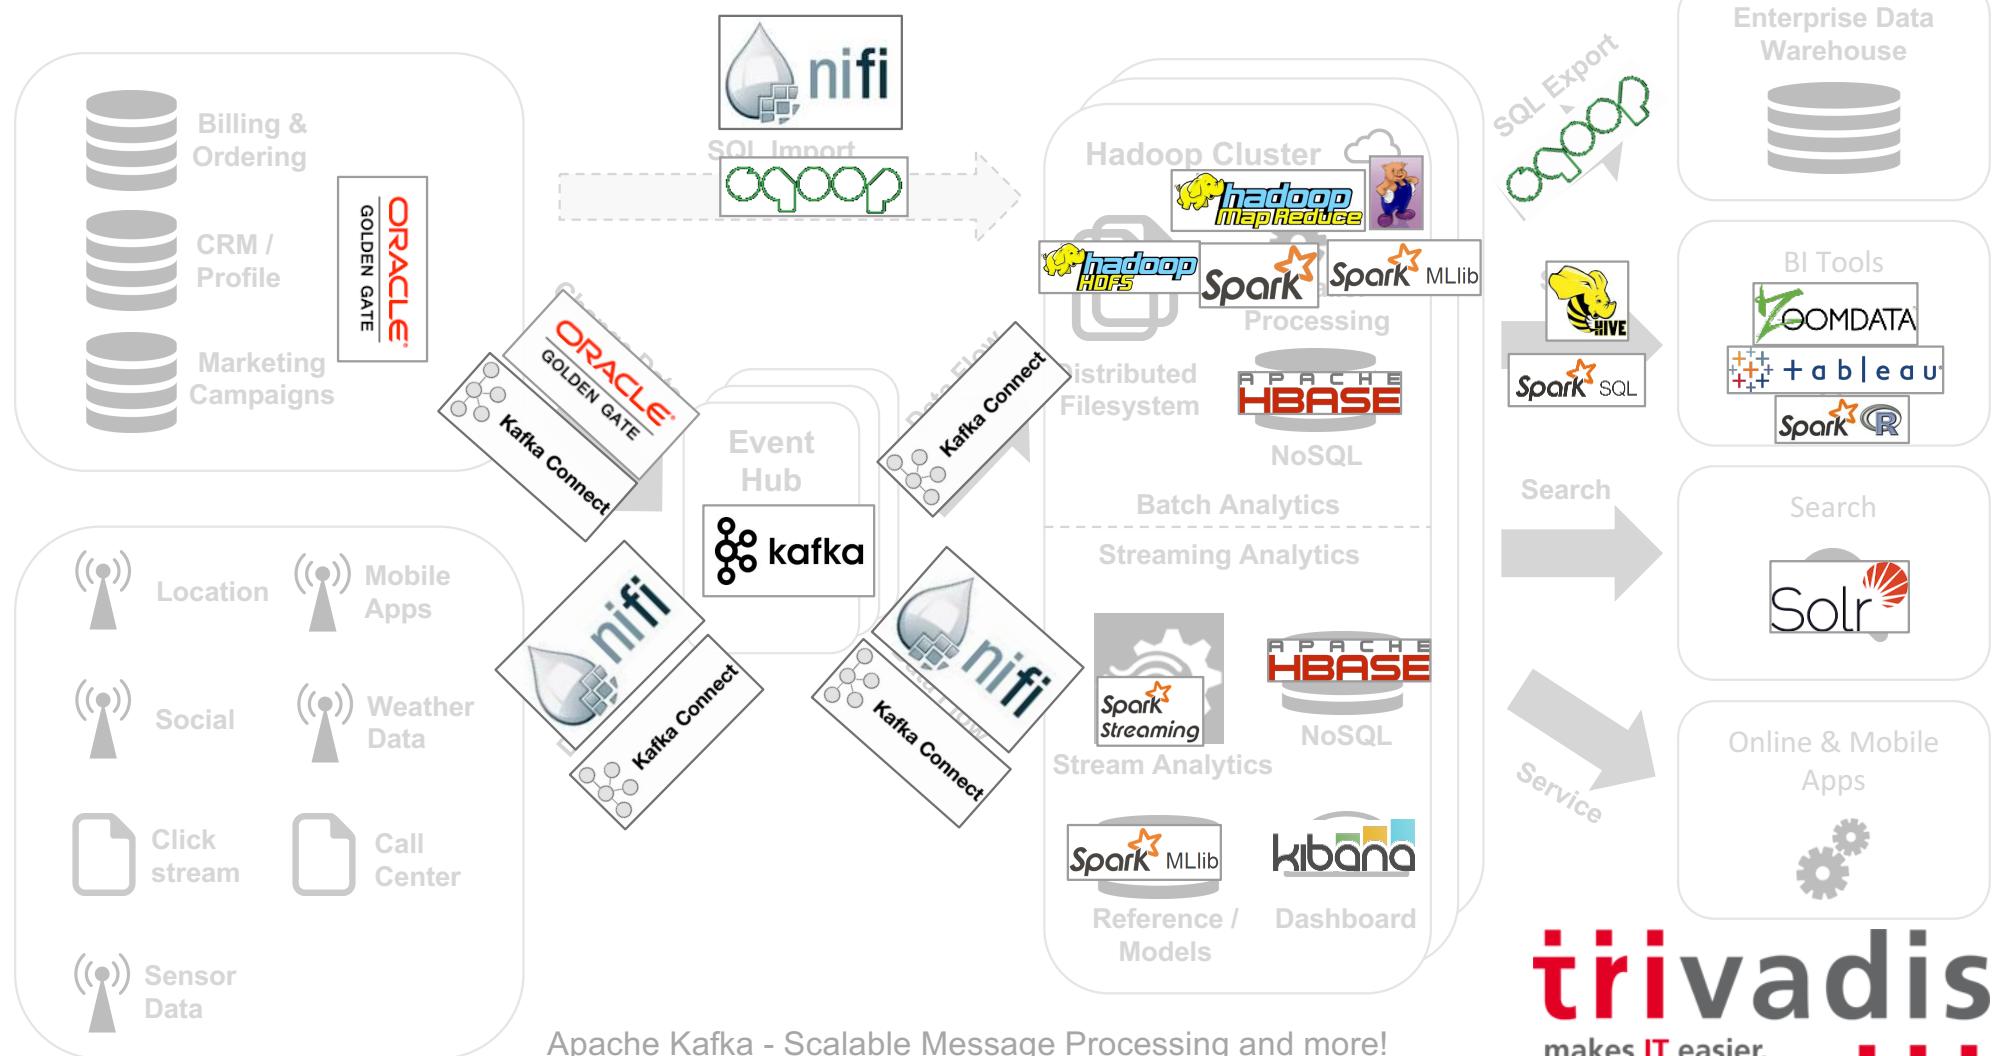
Apache Kafka - Scalable Message Processing and more!

**trivadis**  
makes **IT** easier. 

## ■ Summary

- Kafka can scale to millions of messages per second, and more
- Easy to start in a Proof of Concept (PoC), but more to invest to setup a production environment
- Monitoring is key
- Vibrant community and ecosystem
- Fast paced technology
- Confluent provides distribution and support for Apache Kafka
- Oracle Event Hub Service offers a Kafka Managed Service

# ■ Customer Event Hub – mapping of technologies



**trivadis**  
makes IT easier. ■ ■ ■



**Guido Schmutz**  
Technology Manager

[guido.schmutz@trivadis.com](mailto:guido.schmutz@trivadis.com)



@gschmutz



[guidoschmutz.wordpress.com](http://guidoschmutz.wordpress.com)

Apache Kafka - Scalable Message Processing and more!

**trivadis**  
makes **IT** easier. 